

Linear regression - Logistic regression

Diane Lingrand





2022 - 2023

1 Linear regression

2 Logistic regression


- Context :
 - supervised learning, regression
- Data :
 - m data of dimension :
 - 1 : scalar value : x
 - n : many scalar values : vector : $x = [x_1 x_2 \dots x_n]$
- Goal : we want to predict y value
- Examples :
 - amount of rain from altitude in the Alps
 - price of an apartment from size in m^2
 - vote from age, sex, income, residence location, ...
 - risk of a disease from age, weight, result of blood analysis, ...

Example of linear regression : Ohm law



- Regression : determine value of y with respect to x .
- Linear (dim 1) : the function is a line parameterised by $\theta = [\theta_0 \theta_1]$:


$$y = \theta_0 + x * \theta_1$$



- Learning : determine θ_0 et θ_1
- Regression : compute $y = h_{\theta}(x) = \theta_0 + x * \theta_1$
- Cost function (or error) :


$$J(\theta) = \frac{1}{2m} \sum_{i=0}^{m-1} (h_{\theta}(x^i) - y^i)^2$$

- Cost minimisation $J(\theta)$
 - Let us consider cases where $\theta_0 = 0$: determine θ_1 that minimises $J(\theta)$



- Exhaustive search ?
- Gradient descent

- Find extrema of J : find the zeros of $\frac{dJ}{d\theta_1}$
- Iterative algorithm :
 - pick initial value of θ_1
 - while $J(\theta_1)$ changes (stop when $\frac{dJ}{d\theta_1}(\theta_1) \simeq 0$) :
 - replace θ_1 by $\theta_1 - \alpha \frac{dJ}{d\theta_1}$ ($\alpha > 0$, small)
- α : learning rate
- α has to be chosen carefully :
 - if too small : slow convergence
 - if too big : oscillations
 - \Rightarrow plot $J(\theta)$



- many scalar values : vector $x = [x_1 \dots x_n]$
- linear model : $y = \theta_0 + x_1 \theta_1 + x_2 \theta_2 + \dots + x_n \theta_n$
- Find extrema of J : find the zeros of $\frac{dJ}{d\theta}$
- Iterative algorithm :
 - pick initial value of $\theta = [\theta_0 \dots \theta_n]$
 - while $J(\theta)$ changes (stop when $\frac{dJ}{d\theta}(\theta) \simeq 0$) :
 - replace each θ_i by $\theta_i - \alpha \frac{dJ}{d\theta_i}$ ($\alpha > 0$, small)

- batch gradient descent
 - all training samples for each step
- stochastic gradient descent
 - one training sample for each step (need to shuffle training data)
- mini batch ($b=10$)
 - a set of b training samples for each step

Description at : https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html#sklearn.linear_model.LinearRegression

A toy example :

```
#linear regression object creation
lr = LinearRegression()
#learning
lr.fit(data, labels)
#prediction
pred = lr.predict(data)
#score
score = lr.score(data, labels)
print("train score =", score)
```

- toy dataset : Diabetes dataset¹
 - describe the dataset : nb of data, dimension of data, interval of features. . .
 - split train / test
 - learn on the train dataset
 - metrics on train dataset and on test dataset


1. [https:](https://scikit-learn.org/stable/datasets/toy_dataset.html#diabetes-dataset)

[//scikit-learn.org/stable/datasets/toy_dataset.html#diabetes-dataset](https://scikit-learn.org/stable/datasets/toy_dataset.html#diabetes-dataset)


1 Linear regression


2 Logistic regression

- Supervised learning = learn to predict an output when given an input vector
- We know the class / label y for all training data x
- Logistic regression is a classification method
 - Linear regression leads to values $h_{\theta}(x) \in \mathcal{R}$
 - the idea : values in $[0, 1]$ then thresholding at 0.5
- Also known as logit regression or maximum-entropy classification



- Data separation according to labels (0 or 1).
 - Linear separation : line, plane or hyperplane
 - Non-linear separation : polynomial or gaussian
- Notations :
 - data : $x = [x_1 \ x_2 \dots]$
 - labels : $y \in \{0, 1\}$
 - decision criteria h_θ parameterised by θ
 - $\theta = [\theta_0 \ \theta_1 \ \dots]$






- Decision boundary :
 - line of equation : $\theta_0 + \theta_1 x_1 + \theta_2 x_2 = 0$
 - also written as : $\theta^T x = 0$
- Decision :
 - if $\theta^T x \geq 0$ then $y = 1$
 - if $\theta^T x < 0$ then $y = 0$

$$h_{\theta}(x) = s(\theta^T x)$$

with :


$$s(z) = \frac{1}{1 + e^{-z}}$$



- The decision is :
 - if $h_{\theta}(x) \geq 0.5$ then $y = 1$
 - if $h_{\theta}(x) < 0.5$ then $y = 0$

- data : $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$
- m data
- learning aims at finding θ
- method :
 - error minimisation
 - gradient descent (or other minimisation method)

$$J = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})))$$



For all components θ_j de θ :

$$\theta_j = \theta_j - \alpha \frac{\partial J}{\partial \theta_j}$$

with

$$\frac{\partial J}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

- polynomial model :

$$x = [x_0 x_1 \dots x_n x_0^2 x_1^2 \dots x_0 x_1 x_0 x_2 \dots]$$

- under-fitting
 - add parameters
- over-fitting
 - reduce the number of parameters

- in order to avoid over-fitting
- add $\|\theta\|$ to the cost :

$$J = \|\theta\| - C \frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})))$$

- norm : \mathcal{L}_1 (Lasso), \mathcal{L}_2 (Ridge), Elastic-Net ($\frac{1-\rho}{2}\|\theta\|_2^2 + \rho\|\theta\|_1$)
- many solvers :
 - Coordinate Descent (C++ LIBLINEAR library),
 - Stochastic Average Gradient (SAG) or variant SAGA : good for large dataset
 - Broyden–Fletcher–Goldfarb–Shanno algorithm (small data set only) : family of Newton algorithm
- more details on https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

Description at : https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html

A toy example :

```
#logistic regression object creation
logisticRegr = LogisticRegression()
#learning
logisticRegr.fit(data, labels)
#predicted labels computation
labelsPredicted = logisticRegr.predict(data)
#score computation and display
score = logisticRegr.score(data, labels)
print("train score = ", score)
```

- dataset : digits
 - split train/test
 - learn the classification
 - evaluate the classification (train score, test score)
 - compare with decision trees
- dataset : wine
 - same questions

- Evaluation of classification
 - metrics are useful
 - if the accuracy is 95 %, is it good ? enough ?
- Build a GUI for drawing phone numbers
- Use a classification method to save phone numbers

