

Neural networks

Diane Lingrand



2022 - 2023

Outline

1 A single neuron

2 Perceptron

3 Multi layer perceptron (MLP)

4 Experimentations

1 A single neuron

2 Perceptron

3 Multi layer perceptron (MLP)

4 Experimentations

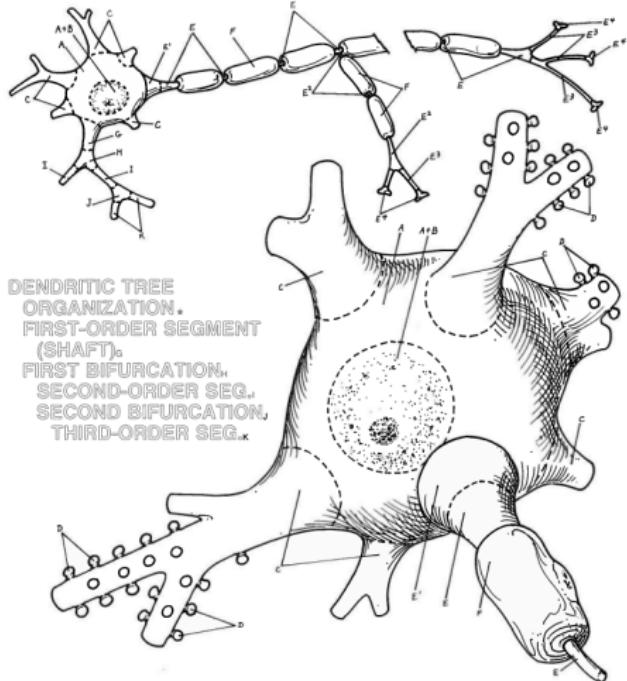
Biological neuron

2-1
THE NEURON

THE NEURON.

NEURON.
CELL BODY.
NUCLEUS.
PROCESSES.
DENDRITE.
DENDRITIC SPINE.

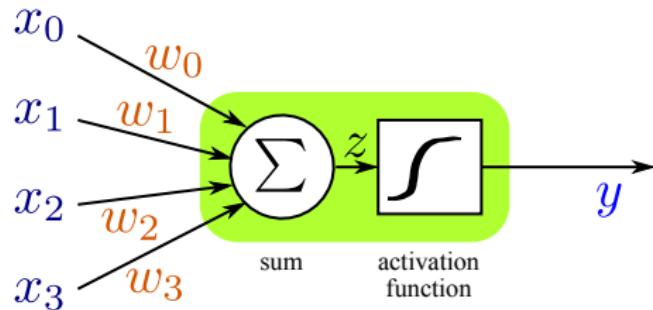
AXON.
AXON HILLOCK.
SHEATH.
AXON COLLATERAL.
TERMINAL BRANCH.
SYNAPTIC TERMINAL.



- dendrites : input.
multiplication with
synaptic weight
- soma or cell body :
summation function
- axon : activation
function and output

The Human Brain Coloring Book
by Diamond et al, 1985

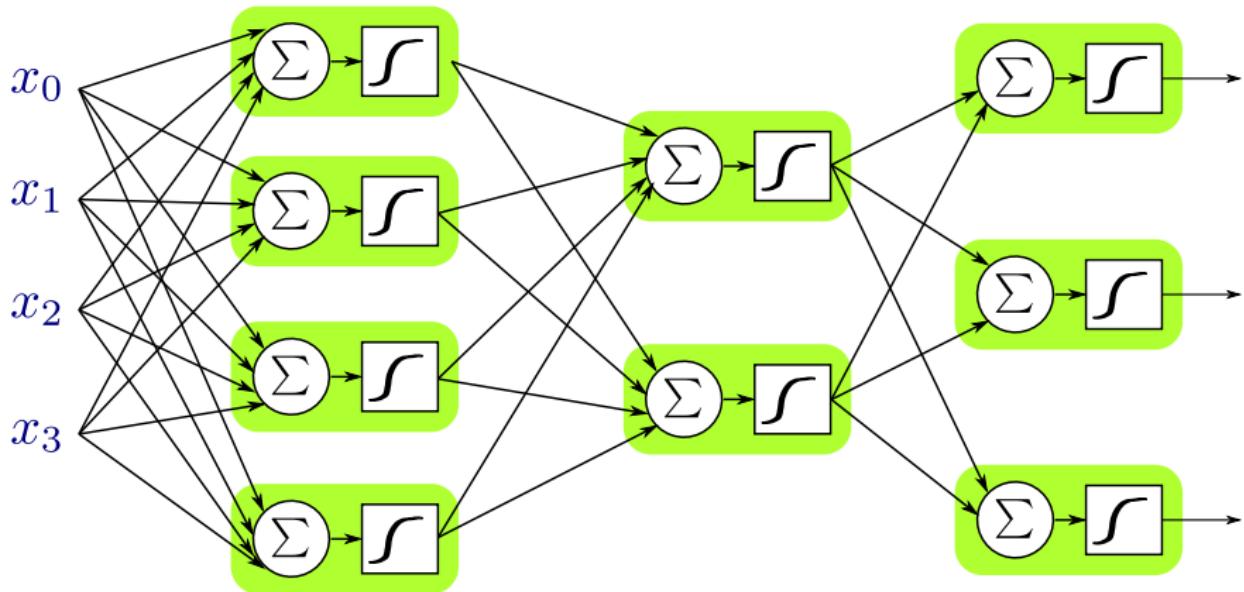
Artificial neuron



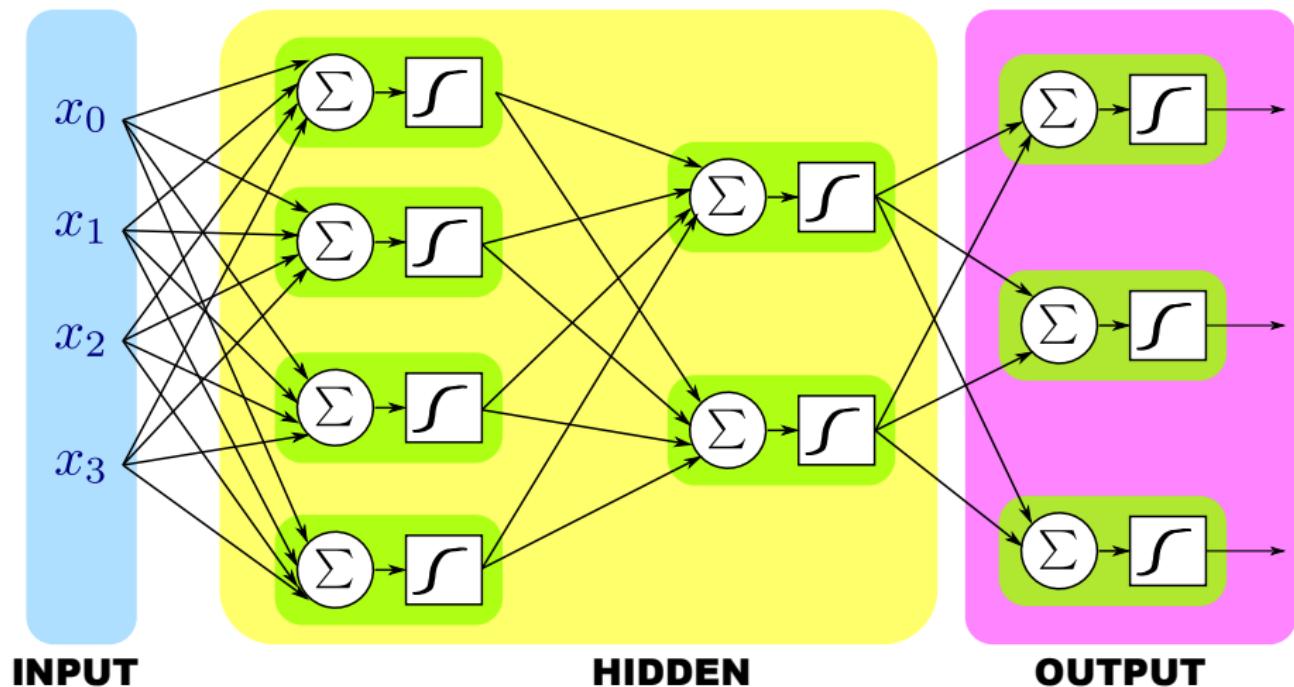
$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

- inputs x_i (also outputs of other neurons)
- associate weights w_i (synaptic modulation)
- bias : special input $x_0 = 1$ with weight w_0
- sum of the weighted inputs (cell body)
- activation function (axon hillock)
- output y

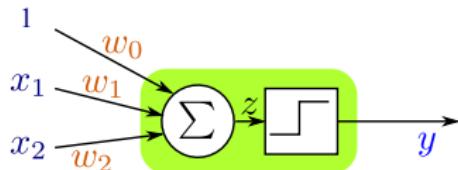
Feed forward neural network



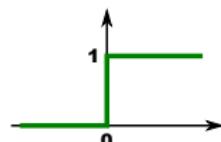
Feed forward neural network



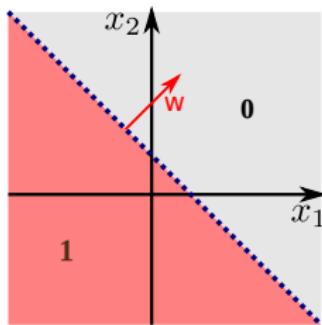
A simple single neuron



- single neuron with 3 inputs :

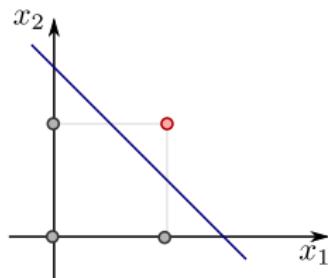


- activation function :
- output : $y = s(z) = s(w_0 + w_1x_1 + w_2x_2)$
- equivalent to divide the 2D plane by a line of equation :
 $w_0 + w_1x_1 + w_2x_2 = 0$ or $w \cdot x = 0$



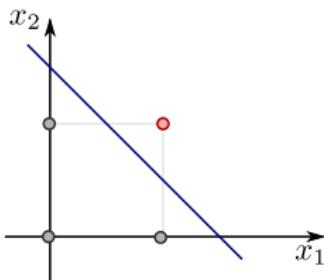
Boolean operation AND

AND	0	1
0	0	0
1	0	1

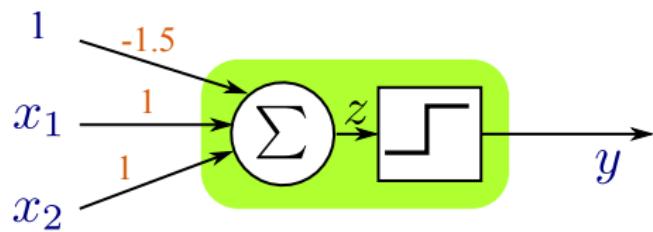


Boolean operation AND

AND	0	1
0	0	0
1	0	1

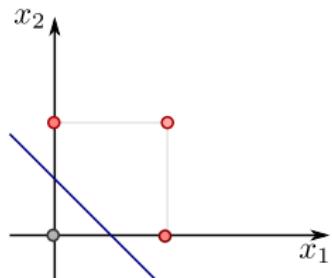


- A solution : $-1.5 + x_1 + x_2 = 0$



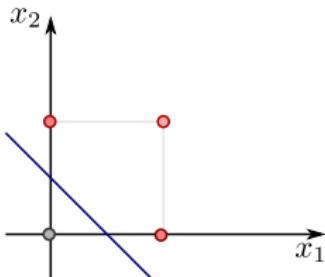
Boolean operation OR

OR	0	1
0	0	1
1	1	1

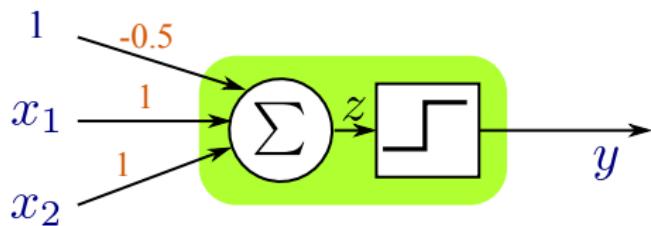


Boolean operation OR

OR	0	1
0	0	1
1	1	1

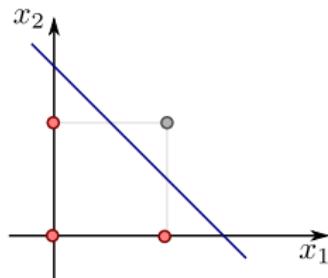


- A solution : $-0.5 + x_1 + x_2 = 0$



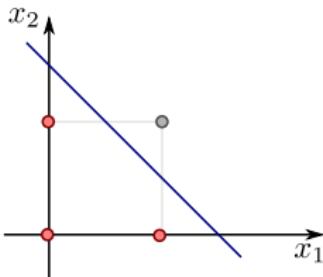
Boolean operation NAND

NAND	0	1
0	1	1
1	1	0

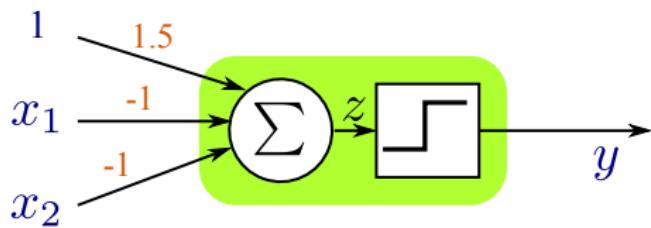


Boolean operation NAND

NAND	0	1
0	1	1
1	1	0

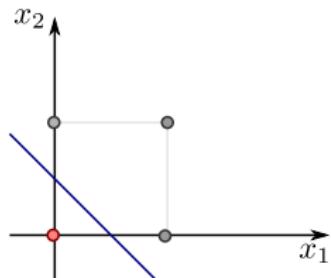


- A solution : $1.5 - x_1 - x_2 = 0$



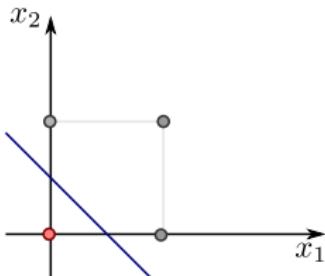
Boolean operation NOR

NOR	0	1
0	0	0
1	1	0

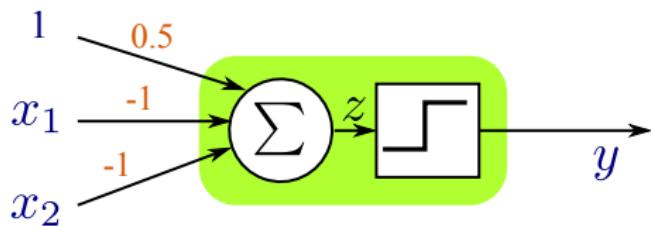


Boolean operation NOR

NOR	0	1
0	0	0
1	1	0

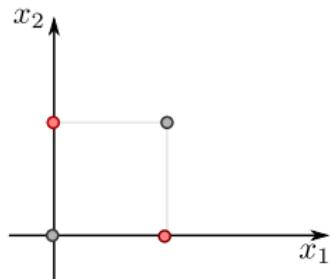


- A solution : $0.5 - x_1 - x_2 = 0$



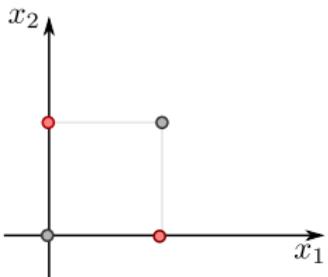
Boolean operation XOR

XOR	0	1
0	0	1
1	1	0



Boolean operation XOR

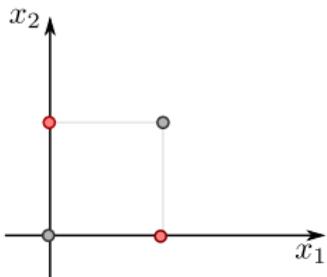
XOR	0	1
0	0	1
1	1	0



- \Rightarrow Impossible with a single neuron !
- { NAND } and { NOR } are functionally complete
- Solution : express XOR using a combination of boolean operators :
 - [A NAND (A NAND B)] NAND [B NAND (A NAND B)]
 - [(A NOR A) NOR (B NOR B)] NOR (A NOR B)

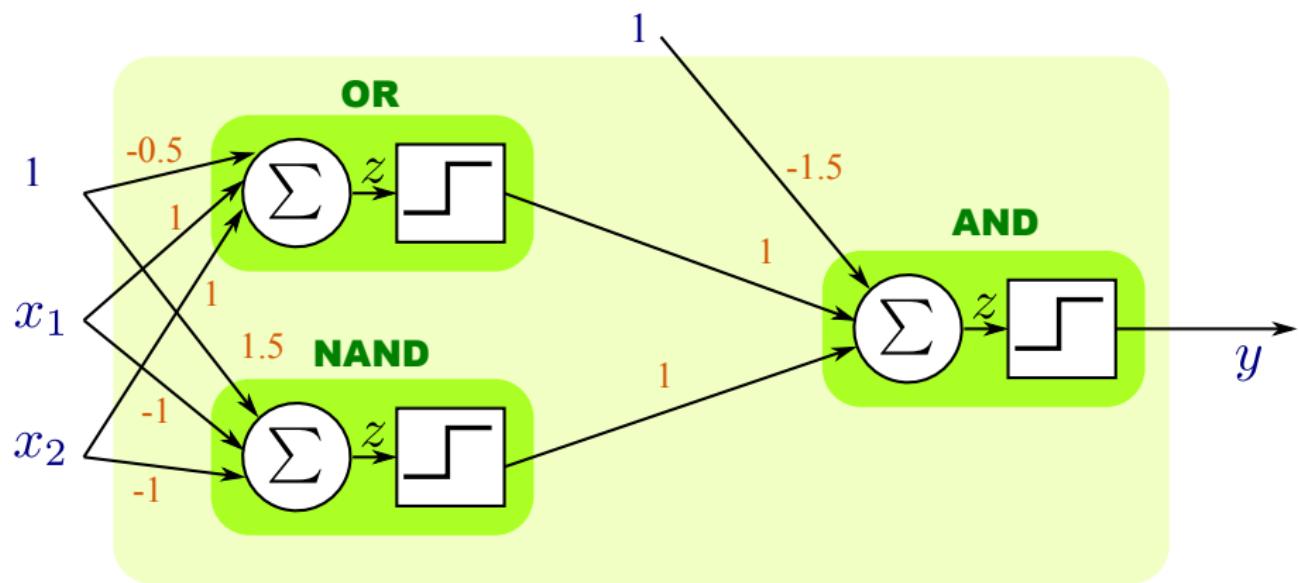
Boolean operation XOR

XOR	0	1
0	0	1
1	1	0



- \Rightarrow Impossible with a single neuron !
- { NAND } and { NOR } are functionally complete
- Solution : express XOR using a combination of boolean operators :
 - [A NAND (A NAND B)] NAND [B NAND (A NAND B)]
 - [(A NOR A) NOR (B NOR B)] NOR (A NOR B)
 - A XOR B = (A NOR B) AND (A OR B)

$$A \text{ XOR } B = (A \text{ NOR } B) \text{ AND } (A \text{ OR } B)$$



- It is proven that (Cybenko and followers) :
 - all boolean functions can be represented using a single hidden layer
 - all bounded continuous functions can be represented using a single hidden layer, with an arbitrary precision
 - all functions can be represented using 2 hidden layers, with an arbitrary precision
- But :
 - How to determine the topology of the neural network ?
 - How many neurons per layer ?
 - How to tune the weights ? (learning phase)

Outline

1 A single neuron

2 Perceptron

3 Multi layer perceptron (MLP)

4 Experimentations

Activation function

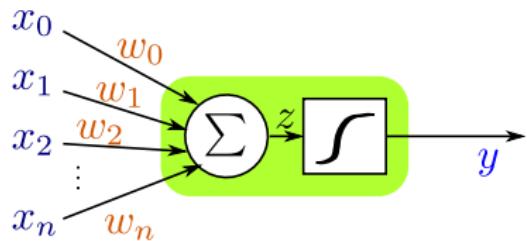
- For the use of gradient, the activation function must be continuous and differentiable.
- Sigmoids (or S-like)

- logistic function $s(z) = \frac{1}{1+e^{-z}}$ and

$$s'(z) = \frac{ds}{dz}(z) = \frac{e^{-z}}{(1+e^{-z})^2} = s(z)(1 - s(z))$$



Back to the single neuron



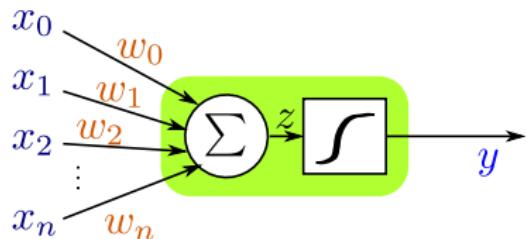
with

$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

For a given data j , the output of the neuron $y(w, x^j)$ should be equal to y^j .

$$E = \frac{1}{2} (y(w, x^j) - y^j)^2$$

Back to the single neuron



with

$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

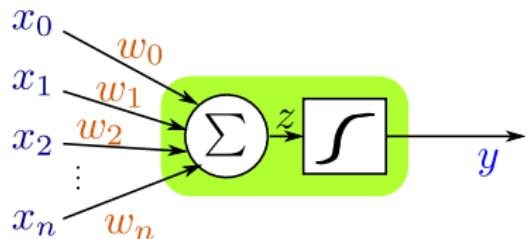
For a given data j , the output of the neuron $y(w, x^j)$ should be equal to y^j .

$$E = \frac{1}{2} (y(w, x^j) - y^j)^2$$

Gradient descent :

$$\nabla_w E = \left[\frac{\partial E}{\partial w_0} \frac{\partial E}{\partial w_1} \cdots \frac{\partial E}{\partial w_n} \right]$$

Back to the single neuron



with

$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

For a given data j , the output of the neuron $y(w, x^j)$ should be equal to y^j .

$$E = \frac{1}{2} (y(w, x^j) - y^j)^2$$

Gradient descent :

$$\nabla_w E = \left[\frac{\partial E}{\partial w_0} \frac{\partial E}{\partial w_1} \cdots \frac{\partial E}{\partial w_n} \right] = (y(w, x^j) - y^j) \nabla_w y(w, x^j)$$

where

$$\nabla_w y(w, x^j) = \left[\frac{\partial y}{\partial w_0} \frac{\partial y}{\partial w_1} \cdots \frac{\partial y}{\partial w_n} \right]$$

Back to the single neuron (2)

We compute :

$$\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j) \text{ where } \nabla_w y(w, x^j) = \left[\frac{\partial y}{\partial w_0} \ \frac{\partial y}{\partial w_1} \ \dots \ \frac{\partial y}{\partial w_n} \right]$$

knowing that :

$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

Back to the single neuron (2)

We compute :

$$\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j) \text{ where } \nabla_w y(w, x^j) = \left[\frac{\partial y}{\partial w_0} \ \frac{\partial y}{\partial w_1} \ \dots \ \frac{\partial y}{\partial w_n} \right]$$

knowing that :

$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

$$\frac{\partial y}{\partial w_i}(w, x^j) =$$

Back to the single neuron (2)

We compute :

$$\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j) \text{ where } \nabla_w y(w, x^j) = \left[\frac{\partial y}{\partial w_0} \frac{\partial y}{\partial w_1} \dots \frac{\partial y}{\partial w_n} \right]$$

knowing that :

$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

$$\frac{\partial y}{\partial w_i}(w, x^j) = \frac{\partial y}{\partial z} \frac{\partial z}{\partial w_i}(w, x^j)$$

Back to the single neuron (2)

We compute :

$$\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j) \text{ where } \nabla_w y(w, x^j) = \left[\frac{\partial y}{\partial w_0} \frac{\partial y}{\partial w_1} \cdots \frac{\partial y}{\partial w_n} \right]$$

knowing that :

$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

$$\frac{\partial y}{\partial w_i}(w, x^j) = \frac{\partial y}{\partial z} \frac{\partial z}{\partial w_i}(w, x^j) = s'(z) \underbrace{\frac{\partial \left(\sum_{k=0}^n w_k x_k^j \right)}{\partial w_i}}_{x_i^j}$$

Back to the single neuron (2)

We compute :

$$\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j) \text{ where } \nabla_w y(w, x^j) = \left[\frac{\partial y}{\partial w_0} \ \frac{\partial y}{\partial w_1} \cdots \frac{\partial y}{\partial w_n} \right]$$

knowing that :

$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

$$\frac{\partial y}{\partial w_i}(w, x^j) = \frac{\partial y}{\partial z} \frac{\partial z}{\partial w_i}(w, x^j) = s'(z) \underbrace{\frac{\partial \left(\sum_{k=0}^n w_k x_k^j \right)}{\partial w_i}}_{x_i^j} = y(1-y)x_i^j$$

with $y = y(w, x^j)$.

Back to the single neuron (2)

We compute :

$$\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j) \text{ where } \nabla_w y(w, x^j) = \left[\frac{\partial y}{\partial w_0} \frac{\partial y}{\partial w_1} \cdots \frac{\partial y}{\partial w_n} \right]$$

knowing that :

$$y = s(z) = s \left(\sum_{i=0}^n w_i x_i \right)$$

$$\frac{\partial y}{\partial w_i}(w, x^j) = \frac{\partial y}{\partial z} \frac{\partial z}{\partial w_i}(w, x^j) = s'(z) \underbrace{\frac{\partial \left(\sum_{k=0}^n w_k x_k^j \right)}{\partial w_i}}_{x_i^j} = y(1-y)x_i^j$$

with $y = y(w, x^j)$. Thus :

$$\nabla_w y = y(1-y)x^j$$

Back to the single neuron (3)

We obtained :

$$\nabla_w E = (y - y^j)y(1 - y)x^j$$

Back to the single neuron (3)

We obtained :

$$\nabla_w E = (y - y^j)y(1 - y)x^j$$

Let us note :

$$\delta = (y - y^j)y(1 - y)$$

Gradient descent rule :

$$w \leftarrow w - \eta \delta x^j$$

Outline

1 A single neuron

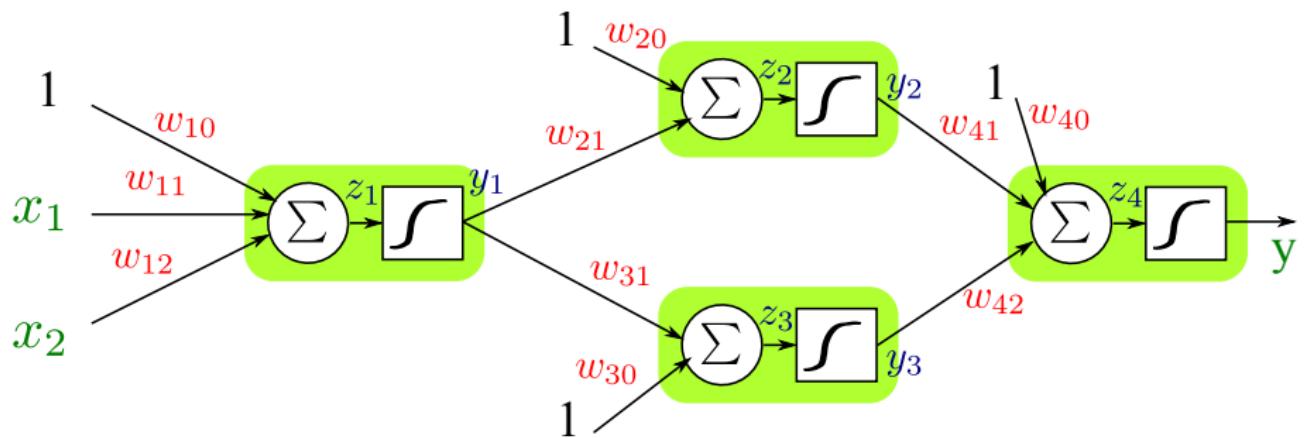
2 Perceptron

3 Multi layer perceptron (MLP)

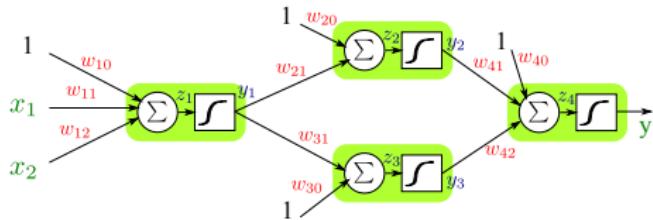
4 Experimentations

- It is proven that (Cybenko and followers) :
 - all boolean functions can be represented using a single hidden layer
 - all bounded continuous functions can be represented using a single hidden layer, with an arbitrary precision
 - all functions can be represented using 2 hidden layers, with an arbitrary precision
- But :
 - How to determine the topology of the neural network ?
 - How many neurons per layer ?
 - How to tune the weights ? (learning phase)

Training a multilayer perceptron

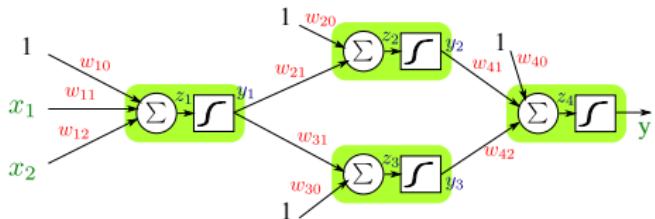


Forward propagation



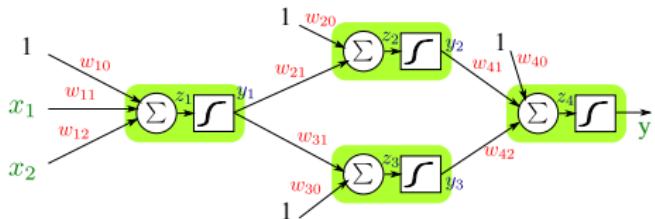
- $y_1 = s(z_1) = s(w_{10} + w_{11}x_1 + w_{12}x_2)$

Forward propagation



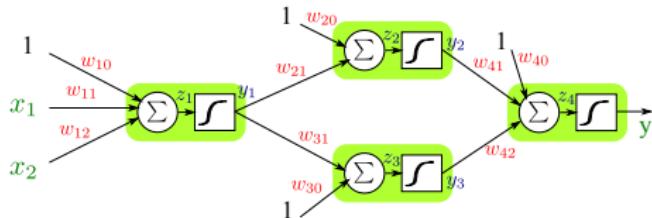
- $y_1 = s(z_1) = s(w_{10} + w_{11}x_1 + w_{12}x_2)$
- $y_2 = s(z_2) = s(w_{20} + w_{21}y_1)$

Forward propagation



- $y_1 = s(z_1) = s(w_{10} + w_{11}x_1 + w_{12}x_2)$
- $y_2 = s(z_2) = s(w_{20} + w_{21}y_1)$
- $y_3 = s(z_3) = s(w_{30} + w_{31}y_1)$

Forward propagation



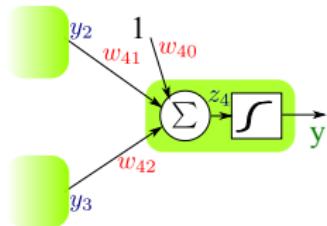
- $y_1 = s(z_1) = s(w_{10} + w_{11}x_1 + w_{12}x_2)$
- $y_2 = s(z_2) = s(w_{20} + w_{21}y_1)$
- $y_3 = s(z_3) = s(w_{30} + w_{31}y_1)$
- $y = s(z_4) = s(w_{40} + w_{41}y_2 + w_{42}y_3)$

Error minimization

- Error : $E = \frac{1}{2}(y(w, x^j) - y^j)^2$ with
 $w = [w_{10} \ w_{11} \ w_{12} \ w_{20} \ w_{21} \ w_{30} \ w_{31} \ w_{40} \ w_{41} \ w_{42}]$ and $x^j = [x_1 \ x_2]$
- Minimization : computation of the gradient
 $\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j)$ where $\nabla_w y = \left[\frac{\partial y}{\partial w_{10}} \ \frac{\partial y}{\partial w_{11}} \cdots \frac{\partial y}{\partial w_{42}} \right]$

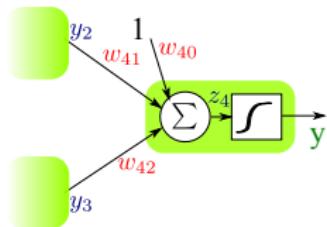
Error minimization

- Error : $E = \frac{1}{2}(y(w, x^j) - y^j)^2$ with
 $w = [w_{10} \ w_{11} \ w_{12} \ w_{20} \ w_{21} \ w_{30} \ w_{31} \ w_{40} \ w_{41} \ w_{42}]$ and $x^j = [x_1 \ x_2]$
- Minimization : computation of the gradient
 $\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j)$ where $\nabla_w y = \left[\frac{\partial y}{\partial w_{10}} \ \frac{\partial y}{\partial w_{11}} \cdots \frac{\partial y}{\partial w_{42}} \right]$
- Output layer :



Error minimization

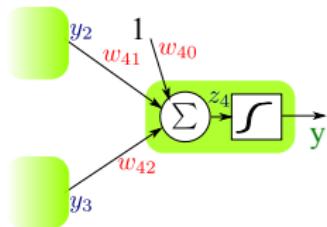
- Error : $E = \frac{1}{2}(y(w, x^j) - y^j)^2$ with
 $w = [w_{10} \ w_{11} \ w_{12} \ w_{20} \ w_{21} \ w_{30} \ w_{31} \ w_{40} \ w_{41} \ w_{42}]$ and $x^j = [x_1 \ x_2]$
- Minimization : computation of the gradient
 $\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j)$ where $\nabla_w y = \left[\frac{\partial y}{\partial w_{10}} \ \frac{\partial y}{\partial w_{11}} \cdots \frac{\partial y}{\partial w_{42}} \right]$
- Output layer :



$$\begin{aligned}\frac{\partial y}{\partial w_{42}} &= \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{42}} = y(1-y)y_3 \\ \frac{\partial y}{\partial w_{41}} &= \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{41}} = y(1-y)y_2 \\ \frac{\partial y}{\partial w_{40}} &= \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{40}} = y(1-y) \\ \Rightarrow \delta_4 &= (y - y^j)y(1-y)\end{aligned}$$

Error minimization

- Error : $E = \frac{1}{2}(y(w, x^j) - y^j)^2$ with
 $w = [w_{10} \ w_{11} \ w_{12} \ w_{20} \ w_{21} \ w_{30} \ w_{31} \ w_{40} \ w_{41} \ w_{42}]$ and $x^j = [x_1 \ x_2]$
- Minimization : computation of the gradient
 $\nabla_w E = (y(w, x^j) - y^j) \nabla_w y(w, x^j)$ where $\nabla_w y = \left[\frac{\partial y}{\partial w_{10}} \ \frac{\partial y}{\partial w_{11}} \cdots \frac{\partial y}{\partial w_{42}} \right]$
- Output layer :

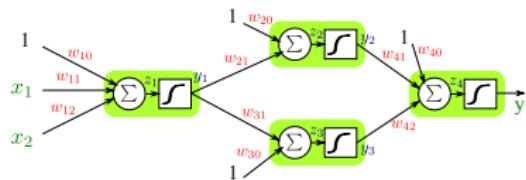


$$\begin{aligned}\frac{\partial y}{\partial w_{42}} &= \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{42}} = y(1-y)y_3 \\ \frac{\partial y}{\partial w_{41}} &= \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{41}} = y(1-y)y_2 \\ \frac{\partial y}{\partial w_{40}} &= \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{40}} = y(1-y) \\ \Rightarrow \delta_4 &= (y - y^j)y(1-y)\end{aligned}$$

- Gradient descent :

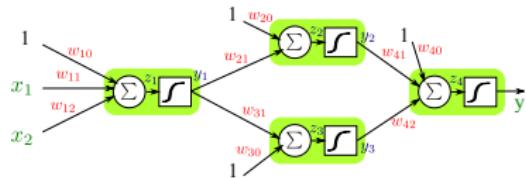
$$\begin{aligned}w_{42} &\leftarrow w_{42} - \eta \delta_4 y_3 \\ w_{41} &\leftarrow w_{41} - \eta \delta_4 y_2 \\ w_{40} &\leftarrow w_{40} - \eta \delta_4\end{aligned}$$

Gradient descent on hidden layers



$$\frac{\partial y}{\partial w_{31}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{31}} \text{ with } \begin{cases} z_4 = w_{40} + w_{41}y_2 + w_{42}y_3 \\ y_3 = s(z_3) = s(w_{30} + w_{31}y_1) \end{cases}$$

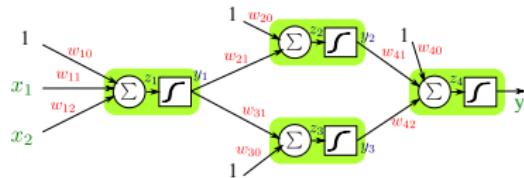
Gradient descent on hidden layers



$$\frac{\partial y}{\partial w_{31}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{31}} \text{ with } \begin{cases} z_4 = w_{40} + w_{41}y_2 + w_{42}y_3 \\ y_3 = s(z_3) = s(w_{30} + w_{31}y_1) \end{cases}$$

- Then $\frac{\partial y}{\partial w_{31}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial y_3} \frac{\partial y_3}{\partial z_3} \frac{\partial z_3}{\partial w_{31}} = y(1-y)w_{42}y_3(1-y_3)y_1$

Gradient descent on hidden layers

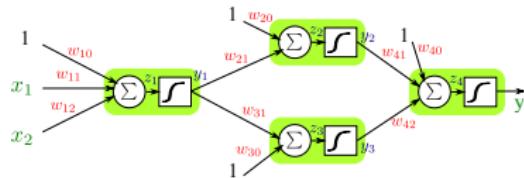


$$\frac{\partial y}{\partial w_{31}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{31}} \text{ with } \begin{cases} z_4 = w_{40} + w_{41}y_2 + w_{42}y_3 \\ y_3 = s(z_3) = s(w_{30} + w_{31}y_1) \end{cases}$$

- Then $\frac{\partial y}{\partial w_{31}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial y_3} \frac{\partial y_3}{\partial z_3} \frac{\partial z_3}{\partial w_{31}} = y(1-y)w_{42}y_3(1-y_3)y_1$
- And, similarly :

$$\frac{\partial y}{\partial w_{30}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{30}} = y(1-y)w_{42}y_3(1-y_3)$$

Gradient descent on hidden layers



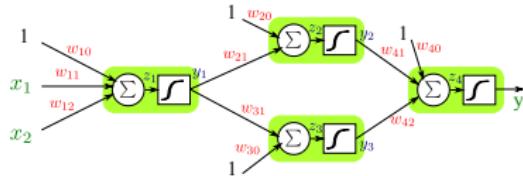
$$\frac{\partial y}{\partial w_{31}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{31}} \text{ with } \begin{cases} z_4 = w_{40} + w_{41}y_2 + w_{42}y_3 \\ y_3 = s(z_3) = s(w_{30} + w_{31}y_1) \end{cases}$$

- Then $\frac{\partial y}{\partial w_{31}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial y_3} \frac{\partial y_3}{\partial z_3} \frac{\partial z_3}{\partial w_{31}} = y(1-y)w_{42}y_3(1-y_3)y_1$
- And, similarly :

$$\frac{\partial y}{\partial w_{30}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{30}} = y(1-y)w_{42}y_3(1-y_3)$$

- Thus : $\delta_3 = \delta_4 w_{42}y_3(1-y_3)$ (recursion !)

Gradient descent on hidden layers



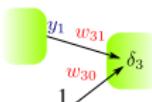
$$\frac{\partial y}{\partial w_{31}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{31}} \text{ with } \begin{cases} z_4 = w_{40} + w_{41}y_2 + w_{42}y_3 \\ y_3 = s(z_3) = s(w_{30} + w_{31}y_1) \end{cases}$$

- Then $\frac{\partial y}{\partial w_{31}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial y_3} \frac{\partial y_3}{\partial z_3} \frac{\partial z_3}{\partial w_{31}} = y(1-y)w_{42}y_3(1-y_3)y_1$
- And, similarly :

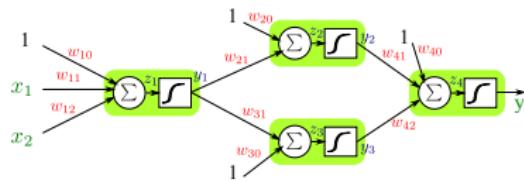
$$\frac{\partial y}{\partial w_{30}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{30}} = y(1-y)w_{42}y_3(1-y_3)$$

- Thus : $\delta_3 = \delta_4 w_{42}y_3(1-y_3)$ (recursion !)

- Gradient descent :
 $w_{30} \leftarrow w_{30} - \eta \delta_3$
 $w_{31} \leftarrow w_{31} - \eta \delta_3 y_1$



Gradient descent on hidden layers (2) (Students have to fill the answers)



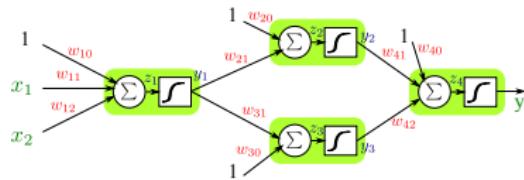
- Similarly :

$$\frac{\partial y}{\partial w_{21}} =$$

- and

$$\frac{\partial y}{\partial w_{20}} =$$

Gradient descent on hidden layers (2) (Students have to fill the answers)



- Similarly :

$$\frac{\partial y}{\partial w_{21}} =$$

- and

$$\frac{\partial y}{\partial w_{20}} =$$

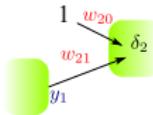
- Thus :

$$\delta_2 = \delta_4 w_{41} y_2 (1 - y_2)$$

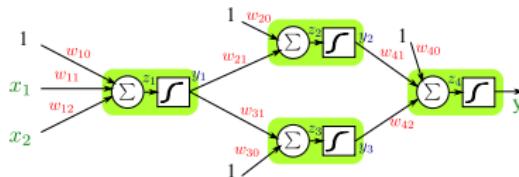
- Gradient descent :

$$w_{20} \leftarrow w_{20} - \eta \delta_2$$

$$w_{21} \leftarrow w_{21} - \eta \delta_2 y_1$$



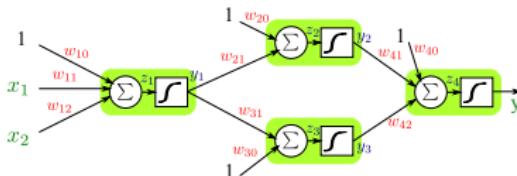
Gradient descent on first hidden layer



$$\frac{\partial y}{\partial w_{12}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{12}}$$

with $z_4 = w_{40} + w_{41}y_2 + w_{42}y_3$

Gradient descent on first hidden layer



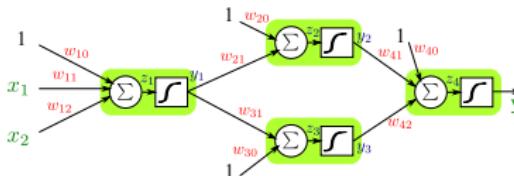
$$\frac{\partial y}{\partial w_{12}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{12}}$$

with $z_4 = w_{40} + w_{41}y_2 + w_{42}y_3$

Thus :

$$\begin{aligned}\frac{\partial y}{\partial w_{12}} &= \frac{\partial y}{\partial z_4} \left(\frac{\partial z_4}{\partial y_2} \frac{\partial y_2}{\partial z_2} \frac{\partial z_2}{\partial y_1} \frac{\partial y_1}{\partial z_1} \frac{\partial z_1}{\partial w_{12}} + \frac{\partial z_4}{\partial y_3} \frac{\partial y_3}{\partial z_3} \frac{\partial z_3}{\partial y_1} \frac{\partial y_1}{\partial z_1} \frac{\partial z_1}{\partial w_{12}} \right) \\&= y(1-y)(w_{41}y_2(1-y_2)w_{21}y_1(1-y_1)x_2 \\&\quad + w_{42}y_3(1-y_3)w_{31}y_1(1-y_1)x_2) \\&= y(1-y)(w_{41}y_2(1-y_2)w_{21} + w_{42}y_3(1-y_3)w_{31}) y_1(1-y_1)x_2\end{aligned}$$

Gradient descent on first hidden layer



$$\frac{\partial y}{\partial w_{12}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{12}}$$

with $z_4 = w_{40} + w_{41}y_2 + w_{42}y_3$

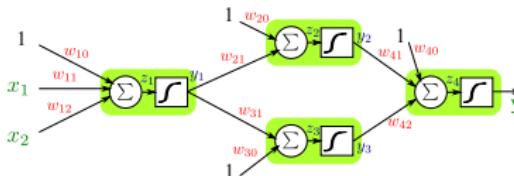
Thus :

$$\begin{aligned}\frac{\partial y}{\partial w_{12}} &= \frac{\partial y}{\partial z_4} \left(\frac{\partial z_4}{\partial y_2} \frac{\partial y_2}{\partial z_2} \frac{\partial z_2}{\partial y_1} \frac{\partial y_1}{\partial z_1} \frac{\partial z_1}{\partial w_{12}} + \frac{\partial z_4}{\partial y_3} \frac{\partial y_3}{\partial z_3} \frac{\partial z_3}{\partial y_1} \frac{\partial y_1}{\partial z_1} \frac{\partial z_1}{\partial w_{12}} \right) \\&= y(1-y)(w_{41}y_2(1-y_2)w_{21}y_1(1-y_1)x_2 \\&\quad + w_{42}y_3(1-y_3)w_{31}y_1(1-y_1)x_2) \\&= y(1-y)(w_{41}y_2(1-y_2)w_{21} + w_{42}y_3(1-y_3)w_{31})y_1(1-y_1)x_2\end{aligned}$$

And :

$$\delta_1 = (\delta_2 w_{21} + \delta_3 w_{31})y_1(1-y_1)$$

Gradient descent on first hidden layer



$$\frac{\partial y}{\partial w_{12}} = \frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial w_{12}}$$

with $z_4 = w_{40} + w_{41}y_2 + w_{42}y_3$

Thus :

$$\begin{aligned}\frac{\partial y}{\partial w_{12}} &= \frac{\partial y}{\partial z_4} \left(\frac{\partial z_4}{\partial y_2} \frac{\partial y_2}{\partial z_2} \frac{\partial z_2}{\partial y_1} \frac{\partial y_1}{\partial z_1} \frac{\partial z_1}{\partial w_{12}} + \frac{\partial z_4}{\partial y_3} \frac{\partial y_3}{\partial z_3} \frac{\partial z_3}{\partial y_1} \frac{\partial y_1}{\partial z_1} \frac{\partial z_1}{\partial w_{12}} \right) \\ &= y(1-y)(w_{41}y_2(1-y_2)w_{21}y_1(1-y_1)x_2 \\ &\quad + w_{42}y_3(1-y_3)w_{31}y_1(1-y_1)x_2) \\ &= y(1-y)(w_{41}y_2(1-y_2)w_{21} + w_{42}y_3(1-y_3)w_{31})y_1(1-y_1)x_2\end{aligned}$$

And :

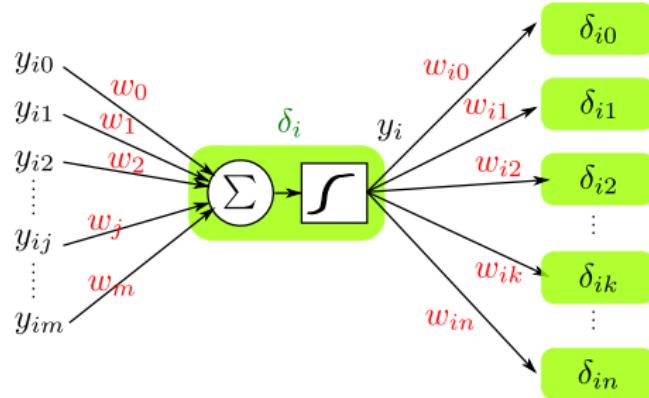
$$\delta_1 = (\delta_2 w_{21} + \delta_3 w_{31})y_1(1-y_1)$$

Gradient descent :

$$\begin{aligned}w_{10} &\leftarrow w_{10} - \eta \delta_1 \\ w_{11} &\leftarrow w_{11} - \eta \delta_1 x_1 \\ w_{12} &\leftarrow w_{12} - \eta \delta_1 x_2\end{aligned}$$

Generalized δ rule

Considering neuron i connected to other neurons :



- Delta rule :

$$\delta_i = y_i(1 - y_i) \sum_{k=0}^n w_{ik} \delta_{ik}$$

- Gradient descent :

$$\forall j \in [0 \ m] \quad w_j \leftarrow w_j - \eta \delta_i y_{ij}$$

Backpropagation Algorithm (*Retropropagation*)

- ① Initialize weights to small random values
- ② Choose a random sample training item, say (x^j, y^j)
- ③ Compute total input z_i and output y_i for each unit (**forward prop**)
- ④ Compute δ_p for output layer $\delta_p = y_p(1 - y_p)(y_p - y_j)$
- ⑤ Compute δ_i for all preceding layers by **backprop rule**
- ⑥ Compute weight change by **descent rule** (repeat for all weights)
 - Note that each expression involves data local to a particular unit, we do not have to look around summing things over the whole network.
 - It is for this reason, simplicity, locality and, therefore, efficiency that backpropagation has become the dominant paradigm for training neural nets.

Input encoding

- All values are real values from [0 1].
- If not, scale the values : $x = \frac{u - u_{\min}}{u_{\max} - u_{\min}}$
- For discrete values, unary encoding form. As for example :

	x_1	x_2	x_3
flour	1	0	0
butter	0	1	0
sugar	0	0	1

Output encoding

- for a **regression** :
 - one output neuron if we need to predict a scalar value
 - n output neurons if the dimension of the vector to predict is n
- for a **binary classification** :
 - one output neuron
 - if $y < 0.5$: class 0
 - else class 1
- for a **multi-class classification** (n classes) :
 - one output neuron : divide $[0 \ 1]$ by n
 - n output neurons
 - each neuron corresponds to probability over classes
 - `tensorflow.keras.utils.to_categorical(yTest, nbClasses)`
 - e.g. : $yTest=3$, $nbClasses = 10$ lead to $[0, 0, 0, 1, 0, 0, 0, 0, 0, 0]^T$
 - `numpy.argmax(model.predict(xTest), axis=1)`
 - e.g. $xTest = [0.01, 0.02, 0.6, 0.05, 0.02, 0.1, 0.1, 0.03, 0.04, 0.03]^T$ leads to 2

- ① Split data set (randomly) into three subsets :
 - Training set : used for adjusting weights
 - Validation set : used to stop training
 - Test set : used to evaluate performance
- ② Pick random, small weights as initial values
- ③ Perform iterative minimization of error over training set
- ④ Stop when error on validation set reaches a minimum (to avoid overfitting)
- ⑤ Repeat training (from step 2) several times (to avoid local minima)
- ⑥ Use best weights to compute error on the test set.

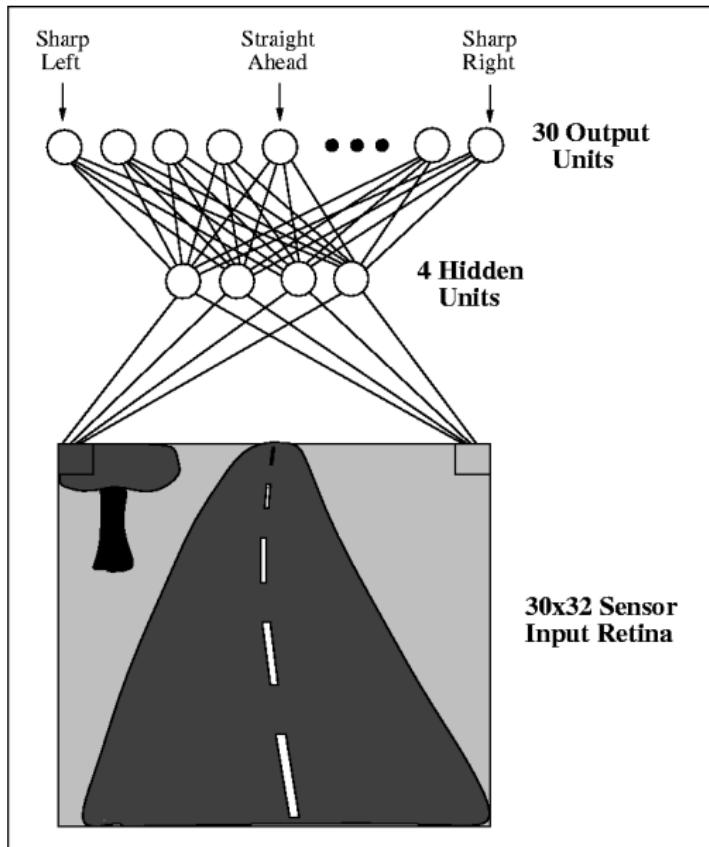
If the data set is too small to be divided into 3 subsets, use cross-validation.

ALVINN : Autonomous Land Vehicle In a Neural Network

- 960 inputs (a 30x32 array derived from the pixels of an image)
- 4 hidden units
- 30 output units (each representing a steering command)

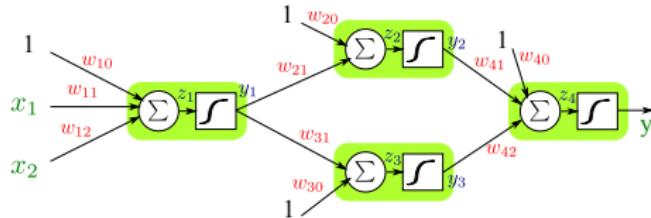
<https://youtu.be/iLP4aPDTBPE>

Dean Pomerleau, CMU, 1989



Exercice

- Here is the network for regression :



- Actual values of weights are :

w_{10}	w_{11}	w_{12}	w_{20}	w_{21}	w_{30}	w_{31}	w_{40}	w_{41}	w_{42}
0.1	0.1	0.1	0.15	0.12	0.09	0.07	0.2	0.2	0.1

- New incoming data is : $x_1 = 0.4$, $x_2 = 0.6$, $y = 0.2$
- Considering $\eta = 0.01$, compute the weights at next step :

w_{10}	w_{11}	w_{12}	w_{20}	w_{21}	w_{30}	w_{31}	w_{40}	w_{41}	w_{42}

- Step 1 : forward propagation. Compute the output of each neuron.
 - $y_1 = s(0.1 + 0.04 + 0.06) = s(0.2) = 0.550$
 - $y_2 = s(0.15 + 0.12 * 0.550) = 0.554$
 - $y_3 = s(0.09 + 0.07 * 0.550) = 0.532$
 - $y = s(0.2 + 0.2 * 0.554 + 0.1 * 0.532) = 0.59$ higher than the expected value of 0.2
- Step 2 : compute the δ for each neuron, backward.
 - $\delta_4 = (0.59 - 0.2) * 0.59 * (1 - 0.59) = 0.094$
 - $\delta_2 = 0.094 * 0.1 * 0.532 * (1 - 0.532) = 0.0047$
 - $\delta_3 = 0.094 * 0.2 * 0.554 * (1 - 0.554) = 0.0023$
 - $\delta_1 = (0.12 * \delta_2 + 0.07 * \delta_3) * 0.55 * (1 - 0.55) = 0.00015$
- Step 3 : update weights

w_{10}	w_{11}	w_{12}	w_{20}	w_{21}
0.099998	0.099999	0.099999	0.14998	0.11999
w_{30}	w_{31}	w_{40}	w_{41}	w_{42}
0.08995	0.06997	0.1991	0.1995	0.0995

new $y_4 = 0.5896$

- Main activation functions :
 - Sigmoid
 - ReLU : Rectified Linear Unit
 - Softmax : transforms output values into values that can be interpreted as probabilities
- Other activation functions :
 - tanh, LeakyReLU, PReLU, ELU, SELU, GELU ...

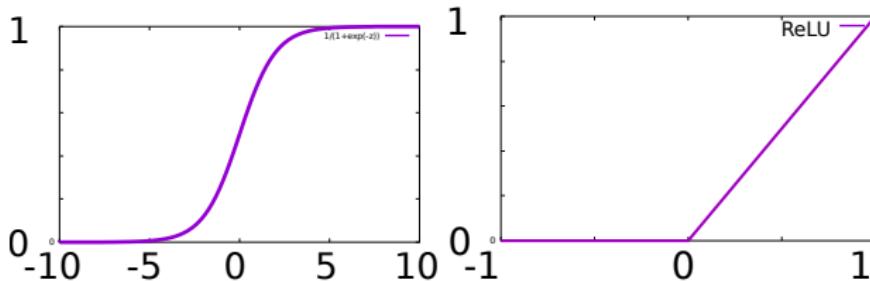
Why ReLU ?

- problem of vanishing gradient
 - remember the delta rule and gradient descent :

$$\delta_i = y_i(1 - y_i) \sum_{k=0}^n w_{ik}\delta_{ik}$$

and

$$\forall j \in [0 \ m] \quad w_j \leftarrow w_j - \eta \delta_i y_{ij}$$



- different solutions :
 - change to activation function to ReLU
 - residual networks
 - batch normalisation

- smooth approximation to the arg max function
 - assign probabilities to indices of having the highest response
 $z = [z_0 z_1 z_2 \dots z_n]$ for each i ,

$$y_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

- every value in $[0, 1]$ and the sum is equal to 1

- classification :
 - binary crossentropy (same as logistic regression) :
 - $-\left(y^{(i)} \log(h_\theta(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)}))\right)$
 - categorical crossentropy :
 - $-\sum_{c=1}^{nbClasses} y^i \log(h_\theta(x^{(i)}))$
- regression :
 - mse, rmse ...
- see https://www.tensorflow.org/api_docs/python/tf/keras/losses for details
 - and this video : <https://www.youtube.com/watch?v=QBbC3Cjsnjg>

Outline

1 A single neuron

2 Perceptron

3 Multi layer perceptron (MLP)

4 Experimentations

How to build a neural network in tensorflow.keras?

- 3 modes :
 - Sequential : the easiest but limited

```
from tensorflow.keras.models import Sequential
model = Sequential()
model.add(Dense(20, input_dim=784, activation='sigmoid'))
```

How to build a neural network in tensorflow.keras ?

- 3 modes :

- Sequential : the easiest but limited

```
from tensorflow.keras.models import Sequential
model = Sequential()
model.add(Dense(20, input_dim=784, activation='sigmoid'))
```

- Functional API : for creating more complex networks (branches, multiple inputs or outputs, DAGs, shared layers)

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense
inputImage = Input(shape=(784,))
myNet = Dense(20, activation='sigmoid')(inputImage)
model = Model(inputImage, myNet)
```

How to build a neural network in tensorflow.keras?

- 3 modes :

- Sequential : the easiest but limited

```
from tensorflow.keras.models import Sequential
model = Sequential()
model.add(Dense(20, input_dim=784, activation='sigmoid'))
```

- Functional API : for creating more complex networks (branches, multiple inputs or outputs, DAGs, shared layers)

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense
inputImage = Input(shape=(784,))
myNet = Dense(20, activation='sigmoid')(inputImage)
model = Model(inputImage, myNet)
```

- Model Subclassing : new from tensorflow 2.0

```
tensorflow.keras.Model class myFirstModel(Model): ...
```

Toy example in python

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation
import numpy as np

X = np.array([[0, 0], [1, 1]])
y = np.array([0,1])

model = Sequential()
model.add(Dense(1, input_dim=2, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])
model.fit(X, y, epochs=100, batch_size=2)
```

Classify handwritten digits using MNIST dataset

Dataset of 60000+10000 grayscale squared images (28x28).



```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation
import numpy as np
from tensorflow.keras.datasets import mnist
```

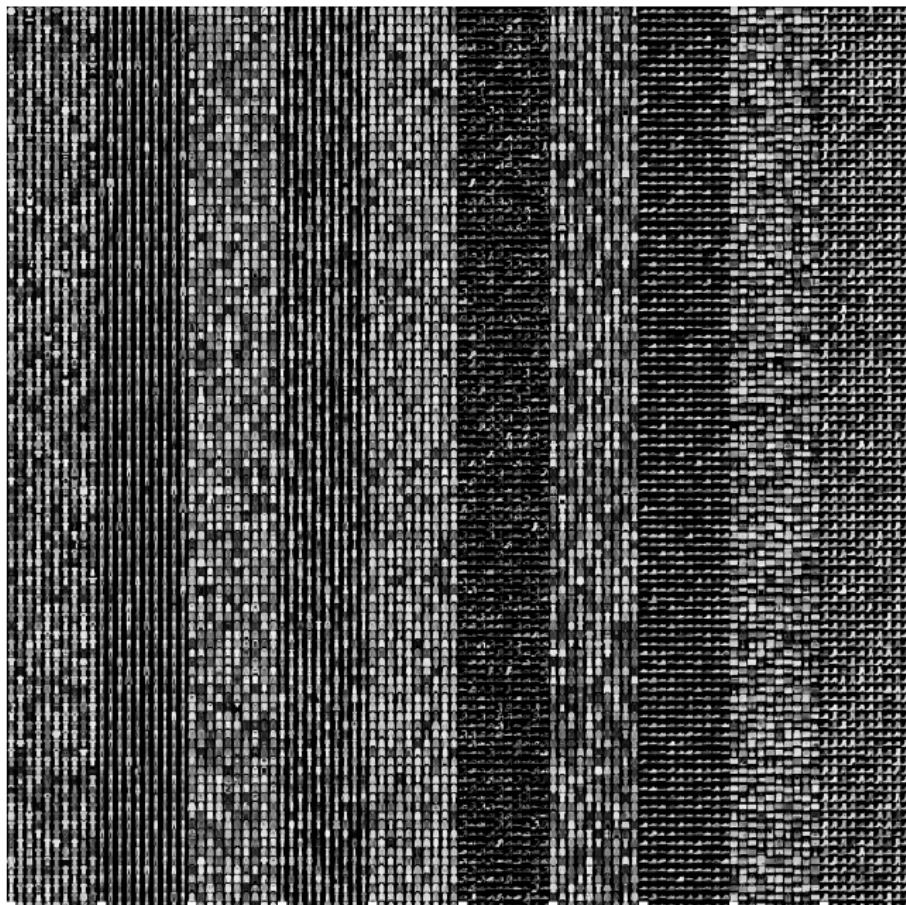
Classify handwritten digits using MNIST dataset : code

```
(xTrain,yTrain),(xTest,yTest) = mnist.load_data()
nbClasses = 10
nbNeuronsHL = 20 # nb of neurons on the hidden layer
xTrain = xTrain.reshape(60000, 784)
xTest = xTest.reshape(10000, 784)
xTrain = xTrain.astype('float32')
xTest = xTest.astype('float32')
xTrain /= 255
xTest /= 255
yTrain = tf.keras.utils.to_categorical(yTrain, nbClasses)
yTest = tf.keras.utils.to_categorical(yTest, nbClasses)
model = Sequential()
model.add(Dense(nbNeuronsHL, input_dim=784, activation='sigmoid'))
model.add(Dense(nbClasses, activation='softmax'))

model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
               metrics=['accuracy'])

model.fit(xTrain, yTrain, epochs=20, batch_size=128)
score = model.evaluate(xTest,yTest)
print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))
```

Fashion MNIST



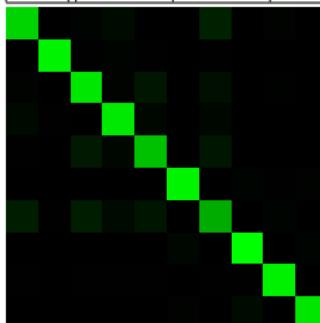
Dataset of
60000+10000
grayscale squared
images (28x28).

Label	Description
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

Tests on fmnist : confusion matrix computed on test data

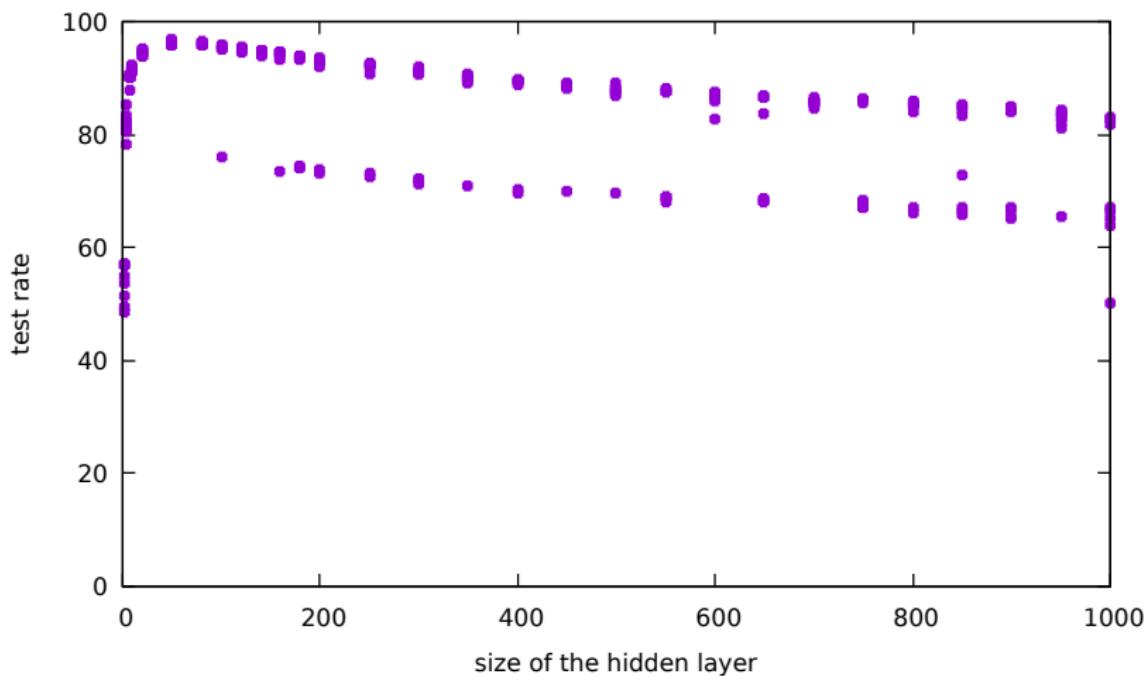
One hidden layer with 20 neurons.

	predicted class									
	0	1	2	3	4	5	6	7	8	9
0	686	3	14	37	3	1	109	1	11	0
1	1	788	6	13	4	0	5	0	1	0
2	11	2	752	11	73	1	49	0	8	0
3	32	6	3	754	21	0	21	0	0	1
4	4	0	81	23	625	0	73	0	2	0
5	0	0	0	0	0	790	0	15	5	11
6	96	4	95	35	71	0	557	2	13	1
7	0	0	0	0	0	20	0	825	0	14
8	5	1	5	5	4	2	6	4	781	0
9	0	0	0	0	0	12	0	38	0	747



One hidden layer : only 10000 data (test) 80%/20%

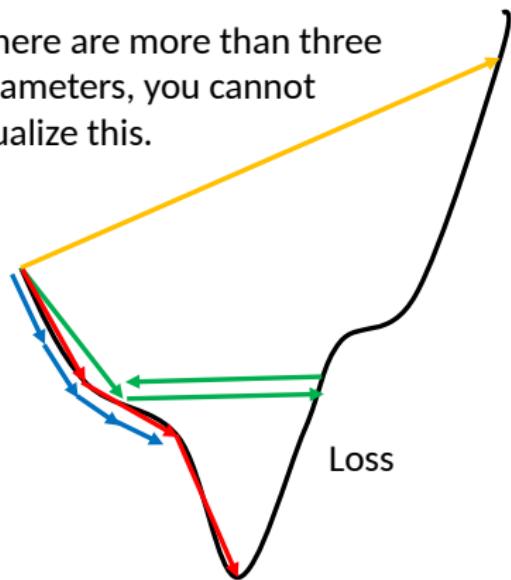
Many trials with random initialisation of weight (modified OpenCV code).



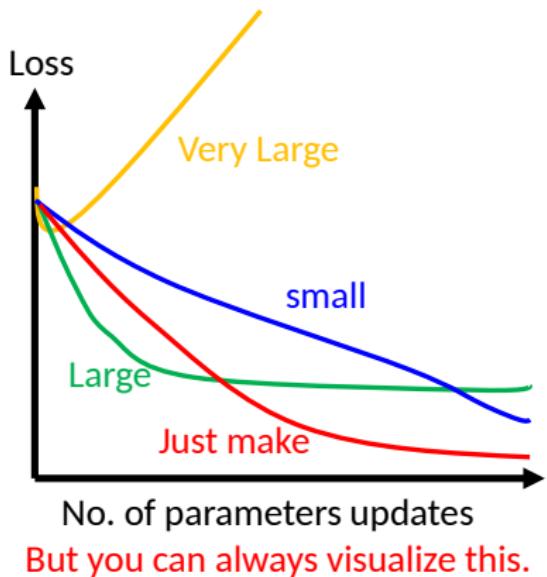
Local minima ...

Learning rate

If there are more than three parameters, you cannot visualize this.

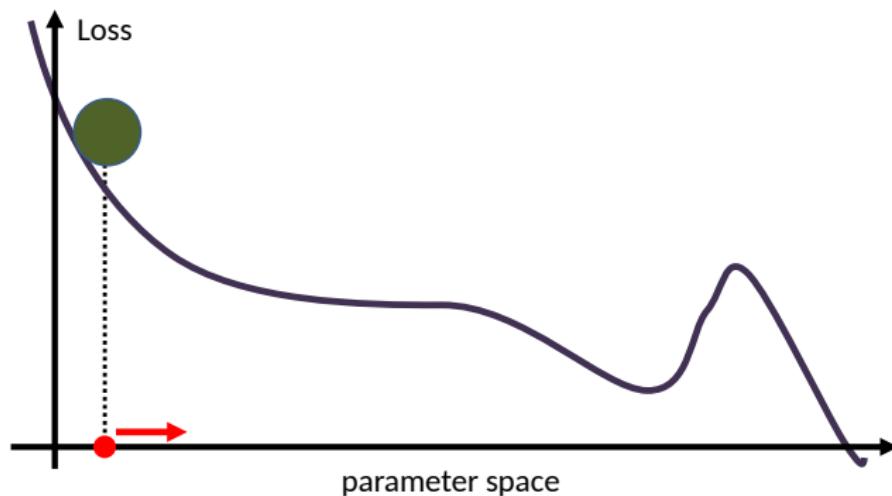


Set the learning rate η carefully

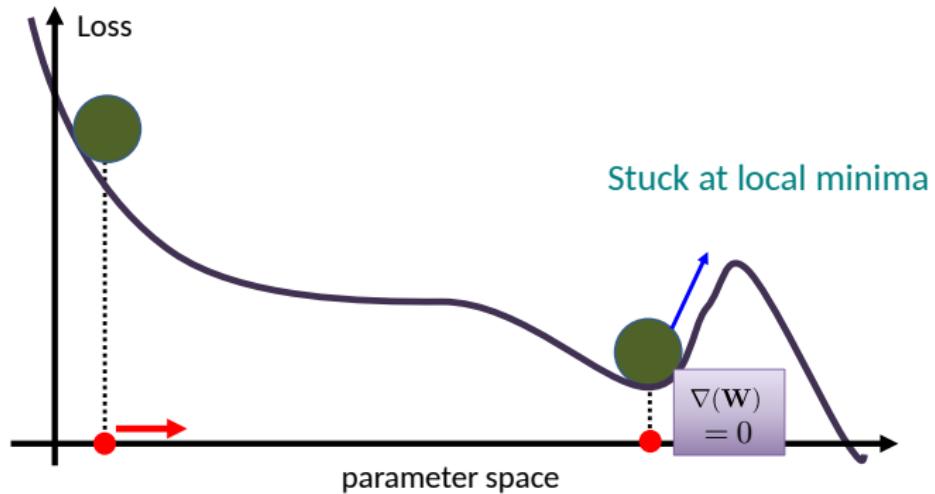


- Only the sign of gradient does matter :
 - if the last 2 gradients have the same sign :
 - good direction, we can fasten the convergence by increasing the learning rate : $\alpha^* = \eta_+$ with $\eta_+ > 1$, e.g. 1.2.
 - if the last 2 gradients have opposite sign :
 - the step was too large : we need to decrease the learning rate : $\alpha^* = \eta_-$ with $\eta_- < 1$, e.g. 0.5.

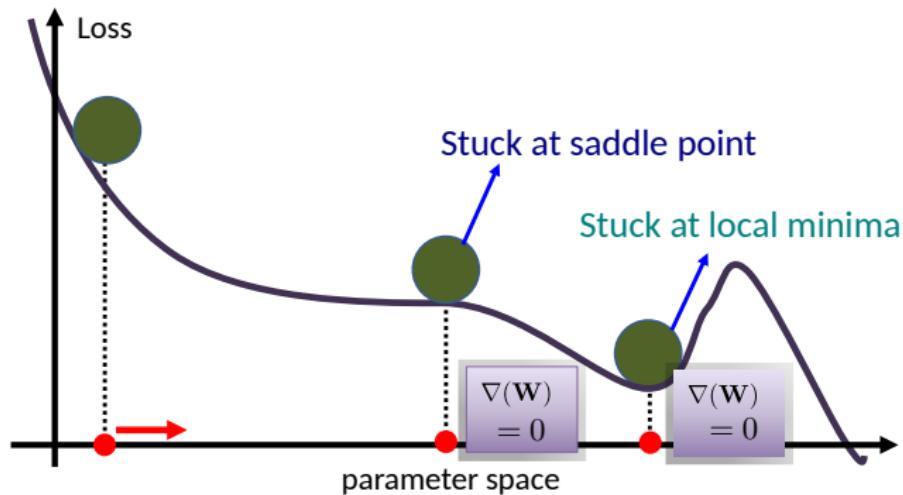
Momentum : avoid local minima



Momentum : avoid local minima

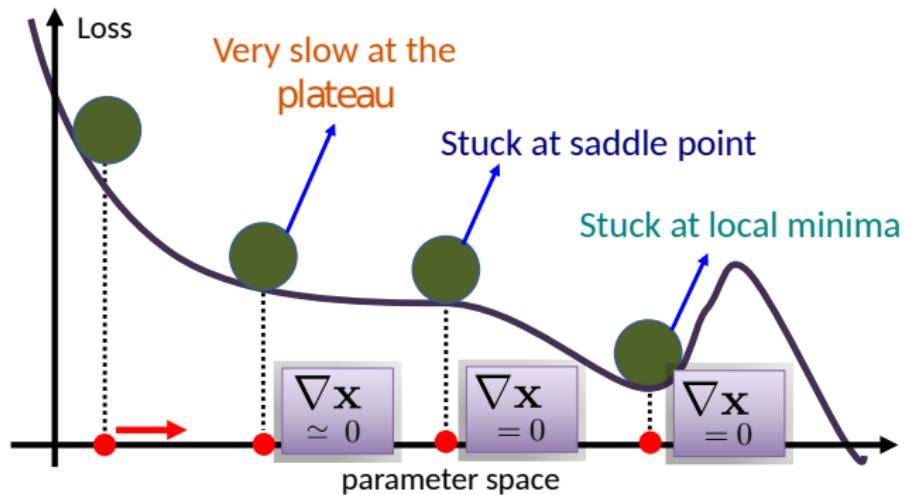


Momentum : avoid local minima

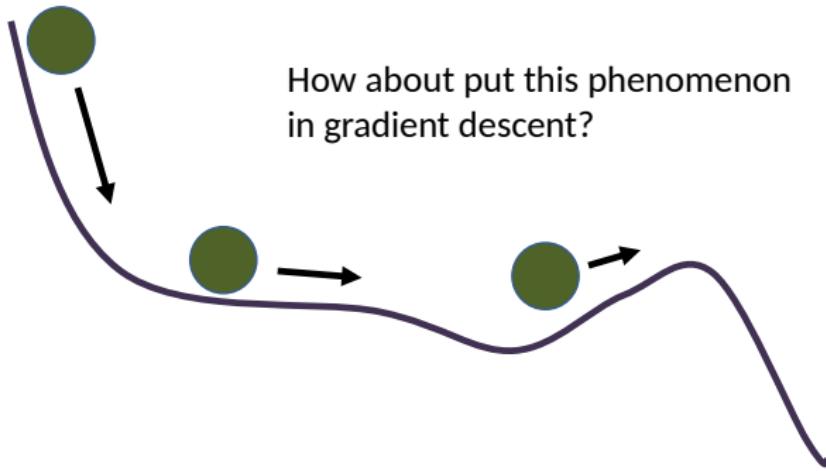


123

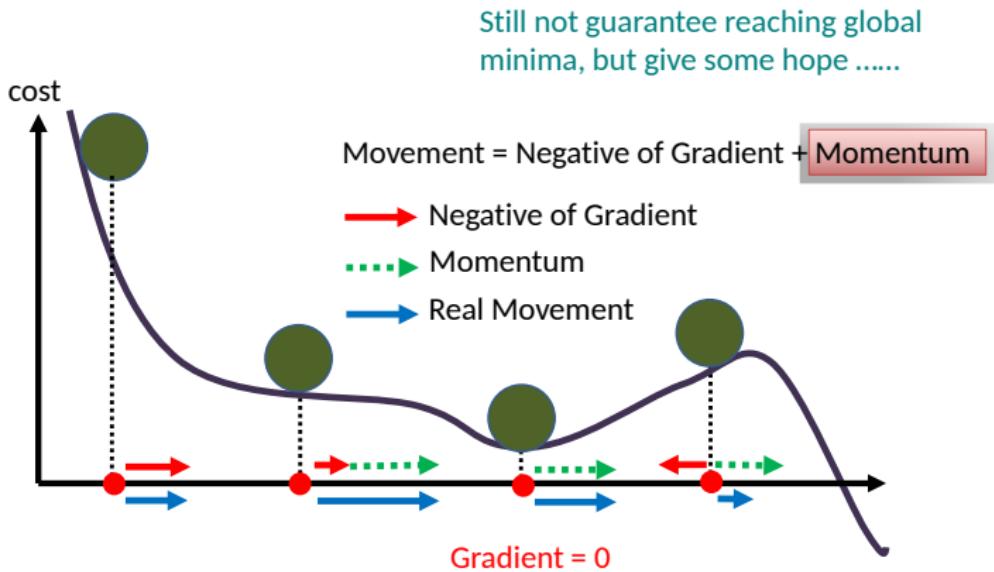
Momentum : avoid local minima



Momentum : avoid local minima



Momentum : avoid local minima



Local minima - Pathological curvatures

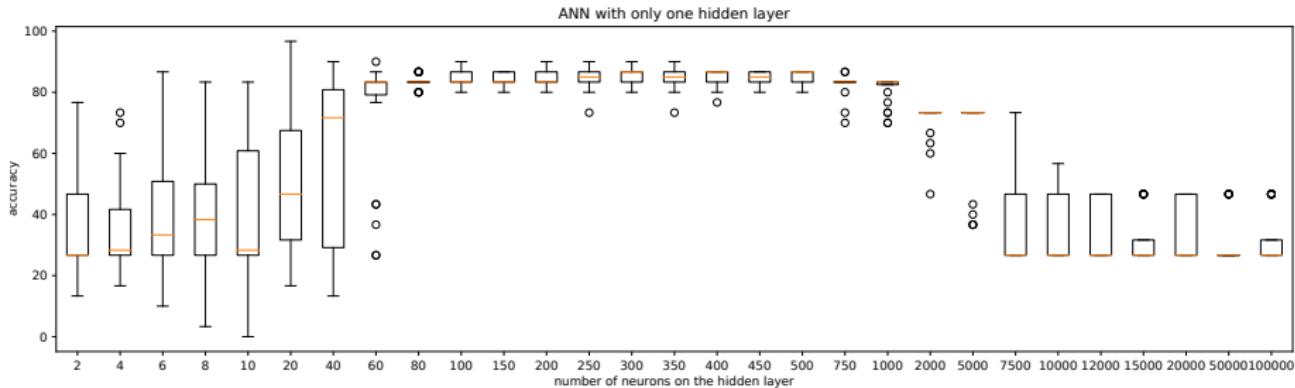
Idea : gradient doesn't tell everything. What about variations of gradient (second derivatives) ?



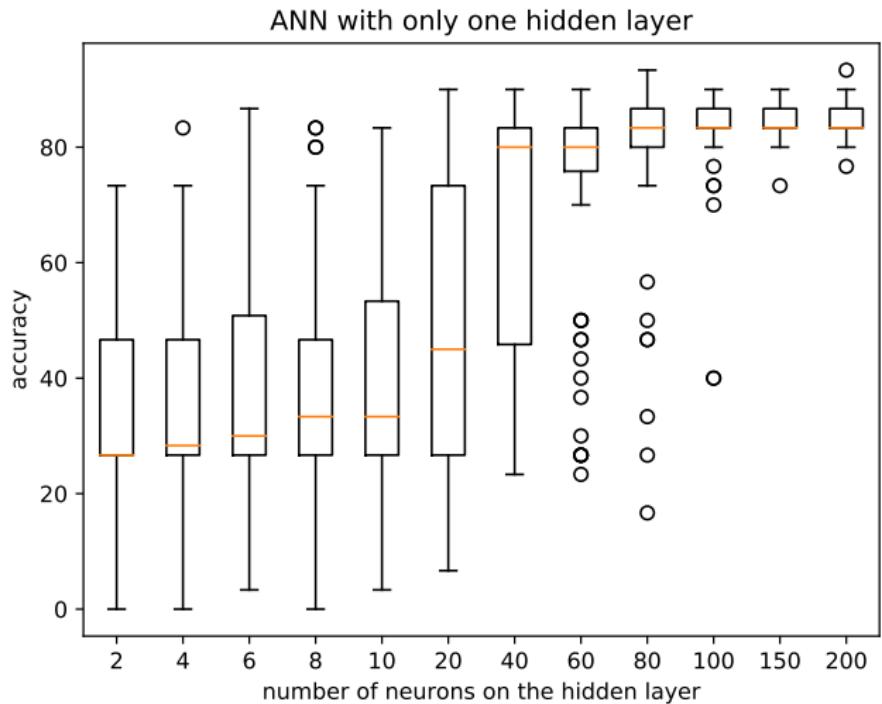
- Momentum
 - accumulates the gradient of the previous steps in order to compute the new weights
 - nice presentation by Andrew Ng
https://www.youtube.com/watch?v=k8fTYJPd3_I
- RMSProp (Root Mean Square Propagation)
 - different learning rates for each weight
 - nice presentation by Andrew Ng
https://www.youtube.com/watch?v=_e-LFe_igno
- AdaM (Adaptive Moment Optimization)
 - combining both Momentum and RMSProp

One hidden layer. Influence of the number of neurons

- Iris dataset (3 classes, 150 data, 4 features).
- ANN using one hidden layer. 24 trials for each set of parameters
- sigmoid activation for hidden layers, softmax for output layer, optimizer RMSprop
- early stopping on validation accuracy, $\min\Delta = 10^{-4}$, patience of 20

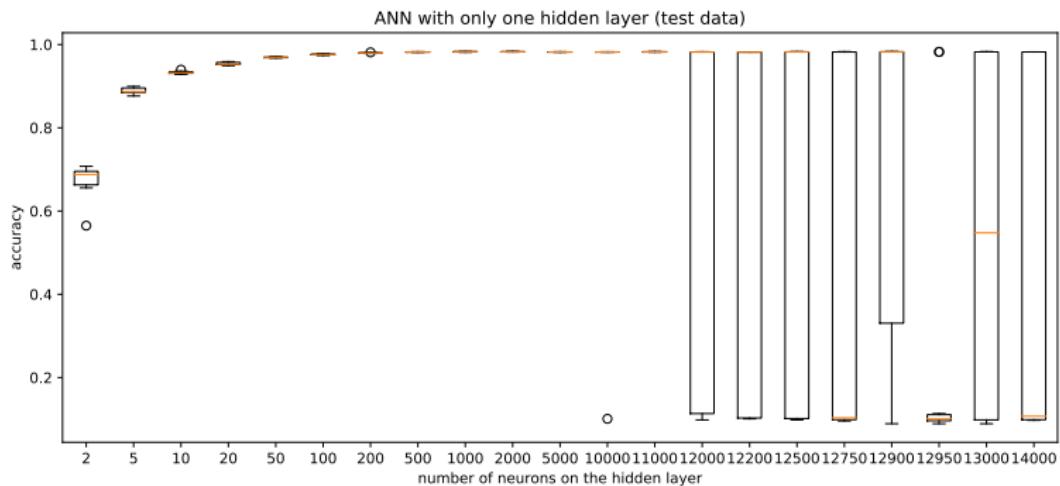


Close-up using 100 trials



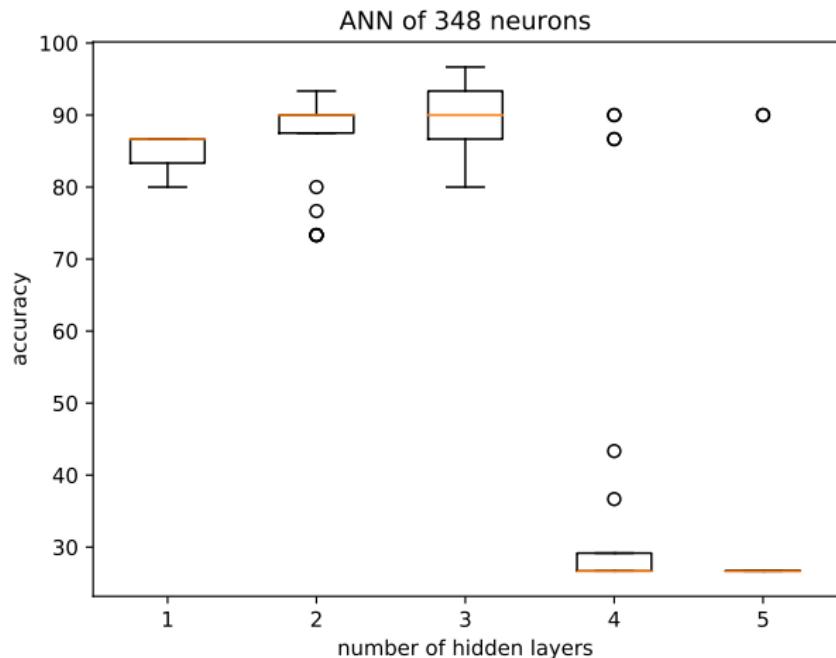
One hidden layer. Influence of the number of neurons

- MNIST data set (10 classes, 70000 data, 784 features)
- ANN using one hidden layer. 10 trials for each set of parameters
- sigmoid activation for hidden layers, softmax for output layer, optimizer RMSprop
- early stopping on validation accuracy, $\min\Delta = 10^{-4}$, patience of 20



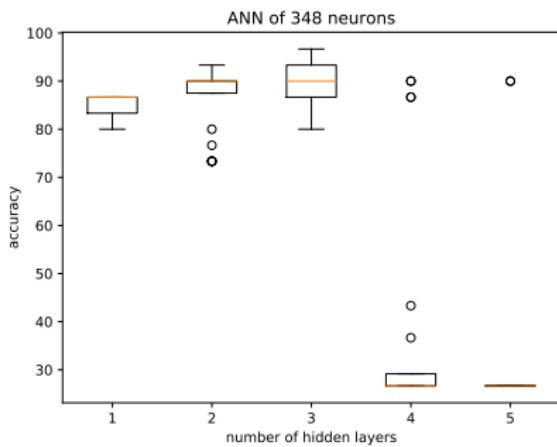
Experimenting the number of hidden layers

- Iris dataset (3 classes, 4 dimensionnal data).
- ANN using 348 neurons on 1 to 5 hidden layers. 24 trials for each set of parameters
- sigmoid activation for hidden layers, softmax for output layer, optimizer RMSprop
- early stopping on validation accuracy, $\text{min}\Delta = 10^{-4}$, patience of 20

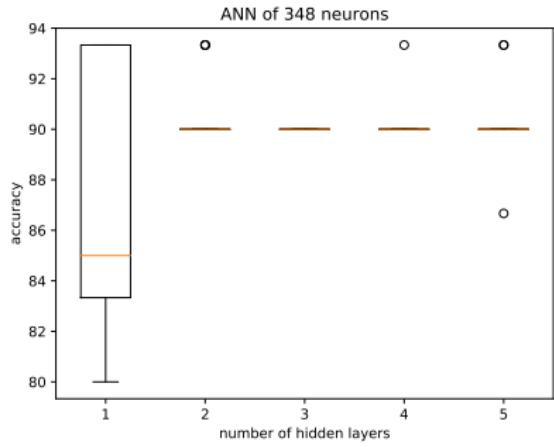


Same experiment comparing sigmoid to ReLu activation

- sigmoid activation



- ReLU activation



Tensorflow playground

<http://playground.tensorflow.org>

- Although Neural Networks kicked off the current phase of interest in Machine Learning, they are extremely problematic
 - Too many parameters (weights, learning rate, momentum, etc)
 - Hard to choose the architecture
 - Very slow to train
 - Easy to get stuck in local minima
- Other methods are of interest